

hdfs fsck命令查看HDFS文件对应的文件块信息(Block)和位置信息(Locations)

在HDFS中，提供了fsck命令，用于检查HDFS上文件和目录的健康状态、获取文件的block信息和位置信息等。

fsck命令必须由HDFS超级用户来执行，普通用户无权限。

```
1. [hadoop@dev ~]$ hdfs fsck
2. Usage: DFSck [-list-corruptfileblocks | [-move | -delete | -openforwrite] [-fi
   les [-blocks [-locations | -racks]]]]
3.         start checking from this path
4.         -move    move corrupted files to /lost+found
5.         -delete  delete corrupted files
6.         -files   print out files being checked
7.         -openforwrite print out files opened for write
8.         -includeSnapshots include snapshot data if the given path indicat
   es a snapshottable directory or there are snapshottable directories under it
9.         -list-corruptfileblocks print out list of missing blocks and files they
   belong to
10.        -blocks  print out block report
11.        -locations print out locations for every block
12.        -racks   print out network topology for data-node locations
```

下面介绍每一个选项的含义及用法。

查看文件中损坏的块 (-list-corruptfileblocks)

```
1. [hadoop@dev ~]$ hdfs fsck /hivedata/warehouse/liuxiaowen.db/lxw_product_names/
   -list-corruptfileblocks
2. The filesystem under path '/hivedata/warehouse/liuxiaowen.db/lxw_product_name
   s/' has 0 CORRUPT files
```

将损坏的文件移动至/lost+found目录 (-move)

```
1. [hadoop@dev ~]$ hdfs fsck /hivedata/warehouse/liuxiaowen.db/lxw_product_names/
   art-00168 -move
2. FSCK started by hadoop (auth:SIMPLE) from /172.16.212.17 for path /hivedata/wa
   rehouse/liuxiaowen.db/lxw_product_names/part-00168 at Thu Aug 13 09:36:35 CST 20
```

```

15
3. .Status: HEALTHY
4. Total size: 13497058 B
5. Total dirs: 0
6. Total files: 1
7. Total symlinks: 0
8. Total blocks (validated): 1 (avg. block size 13497058 B)
9. Minimally replicated blocks: 1 (100.0 %)
10. Over-replicated blocks: 0 (0.0 %)
11. Under-replicated blocks: 0 (0.0 %)
12. Mis-replicated blocks: 0 (0.0 %)
13. Default replication factor: 2
14. Average block replication: 2.0
15. Corrupt blocks: 0
16. Missing replicas: 0 (0.0 %)
17. Number of data-nodes: 15
18. Number of racks: 1
19. FSCK ended at Thu Aug 13 09:36:35 CST 2015 in 1 milliseconds
20.
21.
22. The filesystem under path '/hivedata/warehouse/liuxiaowen.db/lxw_product_names/
    part-00168' is HEALTHY

```

删除损坏的文件 (-delete)

```

1. [hadoop@dev ~]$ hdfs fsck /hivedata/warehouse/liuxiaowen.db/lxw_product_names/p
    art-00168 -delete
2. FSCK started by hadoop (auth:SIMPLE) from /172.16.212.17 for path /hivedata/war
    ehouse/liuxiaowen.db/lxw_product_names/part-00168 at Thu Aug 13 09:37:58 CST 20
    15
3. .Status: HEALTHY
4. Total size: 13497058 B
5. Total dirs: 0
6. Total files: 1
7. Total symlinks: 0
8. Total blocks (validated): 1 (avg. block size 13497058 B)
9. Minimally replicated blocks: 1 (100.0 %)
10. Over-replicated blocks: 0 (0.0 %)
11. Under-replicated blocks: 0 (0.0 %)
12. Mis-replicated blocks: 0 (0.0 %)
13. Default replication factor: 2
14. Average block replication: 2.0
15. Corrupt blocks: 0
16. Missing replicas: 0 (0.0 %)
17. Number of data-nodes: 15
18. Number of racks: 1
19. FSCK ended at Thu Aug 13 09:37:58 CST 2015 in 1 milliseconds
20.
21.

```

```
22. The filesystem under path '/hivedata/warehouse/liuxiaowen.db/lxw_product_names/
part-00168' is HEALTHY
```

检查并列出所有文件状态 (-files)

```
1. [hadoop@dev ~]$ hdfs fsck /hivedata/warehouse/liuxiaowen.db/lxw_product_names/
   -files
2. FSCK started by hadoop (auth:SIMPLE) from /172.16.212.17 for path /hivedata/war
   ouse/liuxiaowen.db/lxw_product_names/ at Thu Aug 13 09:39:38 CST 2015
3. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/ dir
4. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/_SUCCESS 0 bytes, 0
   block(s): OK
5. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00000 13583807 bytes,
   1 block(s): OK
6. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00001 13577427 bytes,
   1 block(s): OK
7. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00002 13588601 bytes,
   1 block(s): OK
8. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00003 13479213 bytes,
   1 block(s): OK
9. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00004 13497012 bytes,
   1 block(s): OK
10. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00005 13557451 bytes,
   1 block(s): OK
11. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00006 13580267 bytes,
   1 block(s): OK
12. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00007 13486035 bytes,
   1 block(s): OK
13. /hivedata/warehouse/liuxiaowen.db/lxw_product_names/part-00008 13481498 bytes,
   1 block(s): OK
14. ...
```

检查并打印正在被打开执行写操作的文件 (-openforwrite)

```
1. [hadoop@dev ~]$ hdfs fsck /hivedata/warehouse/liuxiaowen.db/lxw_product_names/
   -openforwrite
2. FSCK started by hadoop (auth:SIMPLE) from /172.16.212.17 for path /hivedata/war
   ouse/liuxiaowen.db/lxw_product_names/ at Thu Aug 13 09:41:28 CST 2015
3.
   .....
4.
   .....
5. .Status: HEALTHY
6. Total size:      2704782548 B
7. Total dirs:      1
8. Total files:     201
9. Total symlinks:   0
10. Total blocks (validated): 200 (avg. block size 13523912 B)
```

```

11. Minimally replicated blocks: 200 (100.0 %)
12. Over-replicated blocks: 0 (0.0 %)
13. Under-replicated blocks: 0 (0.0 %)
14. Mis-replicated blocks: 0 (0.0 %)
15. Default replication factor: 2
16. Average block replication: 2.0
17. Corrupt blocks: 0
18. Missing replicas: 0 (0.0 %)
19. Number of data-nodes: 15
20. Number of racks: 1
21. FSCK ended at Thu Aug 13 09:41:28 CST 2015 in 10 milliseconds
22.
23. The filesystem under path '/hivedata/warehouse/liuxiaowen.db/lxw_product_name
s/' is HEALTHY

```

打印文件的Block报告 (-blocks)

需要和-files一起使用。

```

1. [hadoop@dev ~]$ hdfs fsck /logs/site/2015-08-08/lxw1234.log -files -blocks
2. FSCK started by hadoop (auth:SIMPLE) from /172.16.212.17 for path /logs/site/20
3. 15-08-08/lxw1234.log at Thu Aug 13 09:45:59 CST 2015
4. /logs/site/2015-08-08/lxw1234.log 7408754725 bytes, 56 block(s): OK
5. 0. BP-1034052771-172.16.212.130-1405595752491:blk_1075892982_2152381 len=134217
6. 728 repl=2
7. 1. BP-1034052771-172.16.212.130-1405595752491:blk_1075892983_2152382 len=134217
8. 728 repl=2
9. 2. BP-1034052771-172.16.212.130-1405595752491:blk_1075892984_2152383 len=134217
10. 728 repl=2
11. 3. BP-1034052771-172.16.212.130-1405595752491:blk_1075892985_2152384 len=134217
12. 728 repl=2
13. 4. BP-1034052771-172.16.212.130-1405595752491:blk_1075892997_2152396 len=134217
14. 728 repl=2
15. 5. BP-1034052771-172.16.212.130-1405595752491:blk_1075892998_2152397 len=134217
16. 728 repl=2
17. 6. BP-1034052771-172.16.212.130-1405595752491:blk_1075892999_2152398 len=134217
18. 728 repl=2
19. 7. BP-1034052771-172.16.212.130-1405595752491:blk_1075893000_2152399 len=134217
20. 728 repl=2
21. 8. BP-1034052771-172.16.212.130-1405595752491:blk_1075893001_2152400 len=134217
22. 728 repl=2
23. 9. BP-1034052771-172.16.212.130-1405595752491:blk_1075893002_2152401 len=134217
24. 728 repl=2
25. 10. BP-1034052771-172.16.212.130-1405595752491:blk_1075893003_2152402 len=13421
26. 7728 repl=2
27. 11. BP-1034052771-172.16.212.130-1405595752491:blk_1075893004_2152403 len=13421
28. 7728 repl=2
29. 12. BP-1034052771-172.16.212.130-1405595752491:blk_1075893005_2152404 len=1342
30. 7728 repl=2

```

```

17. 13. BP-1034052771-172.16.212.130-1405595752491:blk_1075893006_2152405 len=13421
    7728 repl=2
18. 14. BP-1034052771-172.16.212.130-1405595752491:blk_1075893007_2152406 len=13421
    7728 repl=2
19. ...

```

其中，/logs/site/2015-08-08/lxw1234.log 7408754725 bytes, 56 block(s): 表示文件的总大小和block数；

```

0. BP-1034052771-172.16.212.130-1405595752491:blk_1075892982_2152381 len=134217
728 repl=2
1. BP-1034052771-172.16.212.130-1405595752491:blk_1075892983_2152382 len=134217728
repl=2
2. BP-1034052771-172.16.212.130-1405595752491:blk_1075892984_2152383 len=134217728
repl=2

```

前面的0. 1. 2.代表该文件的block索引，56的文件块，就从0-55;

BP-1034052771-172.16.212.130-1405595752491:blk_1075892982_2152381表示block id；

len=134217728 表示该文件块大小；

repl=2 表示该文件块副本数；

打印文件块的位置信息 (-locations)

需要和-files -blocks一起使用。

```

1. [hadoop@dev ~]$ hdfs fsck /logs/site/2015-08-08/lxw1234.log -files -blocks -lo
   cations
2. FSCK started by hadoop (auth:SIMPLE) from /172.16.212.17 for path /logs/site/20
   15-08-08/lxw1234.log at Thu Aug 13 09:45:59 CST 2015
3. /logs/site/2015-08-08/lxw1234.log 7408754725 bytes, 56 block(s): OK
4. 0. BP-1034052771-172.16.212.130-1405595752491:blk_1075892982_2152381 len=134217
   728 repl=2 [172.16.212.139:50010, 172.16.212.135:50010]
5. 1. BP-1034052771-172.16.212.130-1405595752491:blk_1075892983_2152382 len=134217
   728 repl=2 [172.16.212.140:50010, 172.16.212.133:50010]
6. 2. BP-1034052771-172.16.212.130-1405595752491:blk_1075892984_2152383 len=134217
   728 repl=2 [172.16.212.136:50010, 172.16.212.141:50010]
7. 3. BP-1034052771-172.16.212.130-1405595752491:blk_1075892985_2152384 len=134217
   728 repl=2 [172.16.212.133:50010, 172.16.212.135:50010]
8. 4. BP-1034052771-172.16.212.130-1405595752491:blk_1075892997_2152396 len=134217
   728 repl=2 [172.16.212.142:50010, 172.16.212.139:50010]
9. 5. BP-1034052771-172.16.212.130-1405595752491:blk_1075892998_2152397 len=134217
   728 repl=2 [172.16.212.133:50010, 172.16.212.139:50010]
10. 6. BP-1034052771-172.16.212.130-1405595752491:blk_1075892999_2152398 len=134217
    728 repl=2 [172.16.212.141:50010, 172.16.212.135:50010]

```

```
11. 7. BP-1034052771-172.16.212.130-1405595752491:blk_1075893000_2152399 len=134217
    728 repl=2 [172.16.212.144:50010, 172.16.212.142:50010]
12. 8. BP-1034052771-172.16.212.130-1405595752491:blk_1075893001_2152400 len=134217
    728 repl=2 [172.16.212.133:50010, 172.16.212.138:50010]
13. 9. BP-1034052771-172.16.212.130-1405595752491:blk_1075893002_2152401 len=134217
    728 repl=2 [172.16.212.140:50010, 172.16.212.134:50010]
14. ...
```

和打印出的文件块信息相比，多了一个文件块的位置信息：[172.16.212.139:50010, 172.16.212.135:50010]

打印文件块位置所在的机架信息（-racks）

```
1. [hadoop@dev ~]$ hdfs fsck /logs/site/2015-08-08/lxw1234.log -files -blocks -lo
   cations -racks
2. FSCK started by hadoop (auth:SIMPLE) from /172.16.212.17 for path /logs/site/20
   15-08-08/lxw1234.log at Thu Aug 13 09:45:59 CST 2015
3. /logs/site/2015-08-08/lxw1234.log 7408754725 bytes, 56 block(s): OK
4. 0. BP-1034052771-172.16.212.130-1405595752491:blk_1075892982_2152381 len=134217
   728 repl=2 [/default-rack/172.16.212.139:50010, /default-rack/172.16.212.135:50
   010]
5. 1. BP-1034052771-172.16.212.130-1405595752491:blk_1075892983_2152382 len=134217
   728 repl=2 [/default-rack/172.16.212.140:50010, /default-rack/172.16.212.133:50
   010]
6. 2. BP-1034052771-172.16.212.130-1405595752491:blk_1075892984_2152383 len=134217
   728 repl=2 [/default-rack/172.16.212.136:50010, /default-rack/172.16.212.141:50
   010]
7. 3. BP-1034052771-172.16.212.130-1405595752491:blk_1075892985_2152384 len=134217
   728 repl=2 [/default-rack/172.16.212.133:50010, /default-rack/172.16.212.135:50
   010]
8. 4. BP-1034052771-172.16.212.130-1405595752491:blk_1075892997_2152396 len=134217
   728 repl=2 [/default-rack/172.16.212.142:50010, /default-rack/172.16.212.139:50
   010]
9. 5. BP-1034052771-172.16.212.130-1405595752491:blk_1075892998_2152397 len=134217
   728 repl=2 [/default-rack/172.16.212.133:50010, /default-rack/172.16.212.139:50
   010]
10. ...
```

和前面打印出的信息相比，多了机架信息：[/default-rack/172.16.212.139:50010, /default-rack/172.16.212.135:50010]

转载请注明：lxw的大数据田地 » [hdfs fsck命令查看HDFS文件对应的文件块信息\(Block\)和位置信息\(Locations\)](#)